

Intelligent agents: individual development assignment

Guilherme Amorim

April 2025




Introduction



Intelligent agents¹

- Sense environment
- Take action autonomously to achieve specific goals



Potential value^{2,3}

- Efficiency gains
(automation, complex decision-making)
- Foster innovation
(new business models)



Use-cases⁴⁻⁶

- Healthcare
(precision diagnosis, tailored treatment recommendations)
- eCommerce
(personalized shopping experiences)
- Finance
(fraud detection, algorithmic trading)

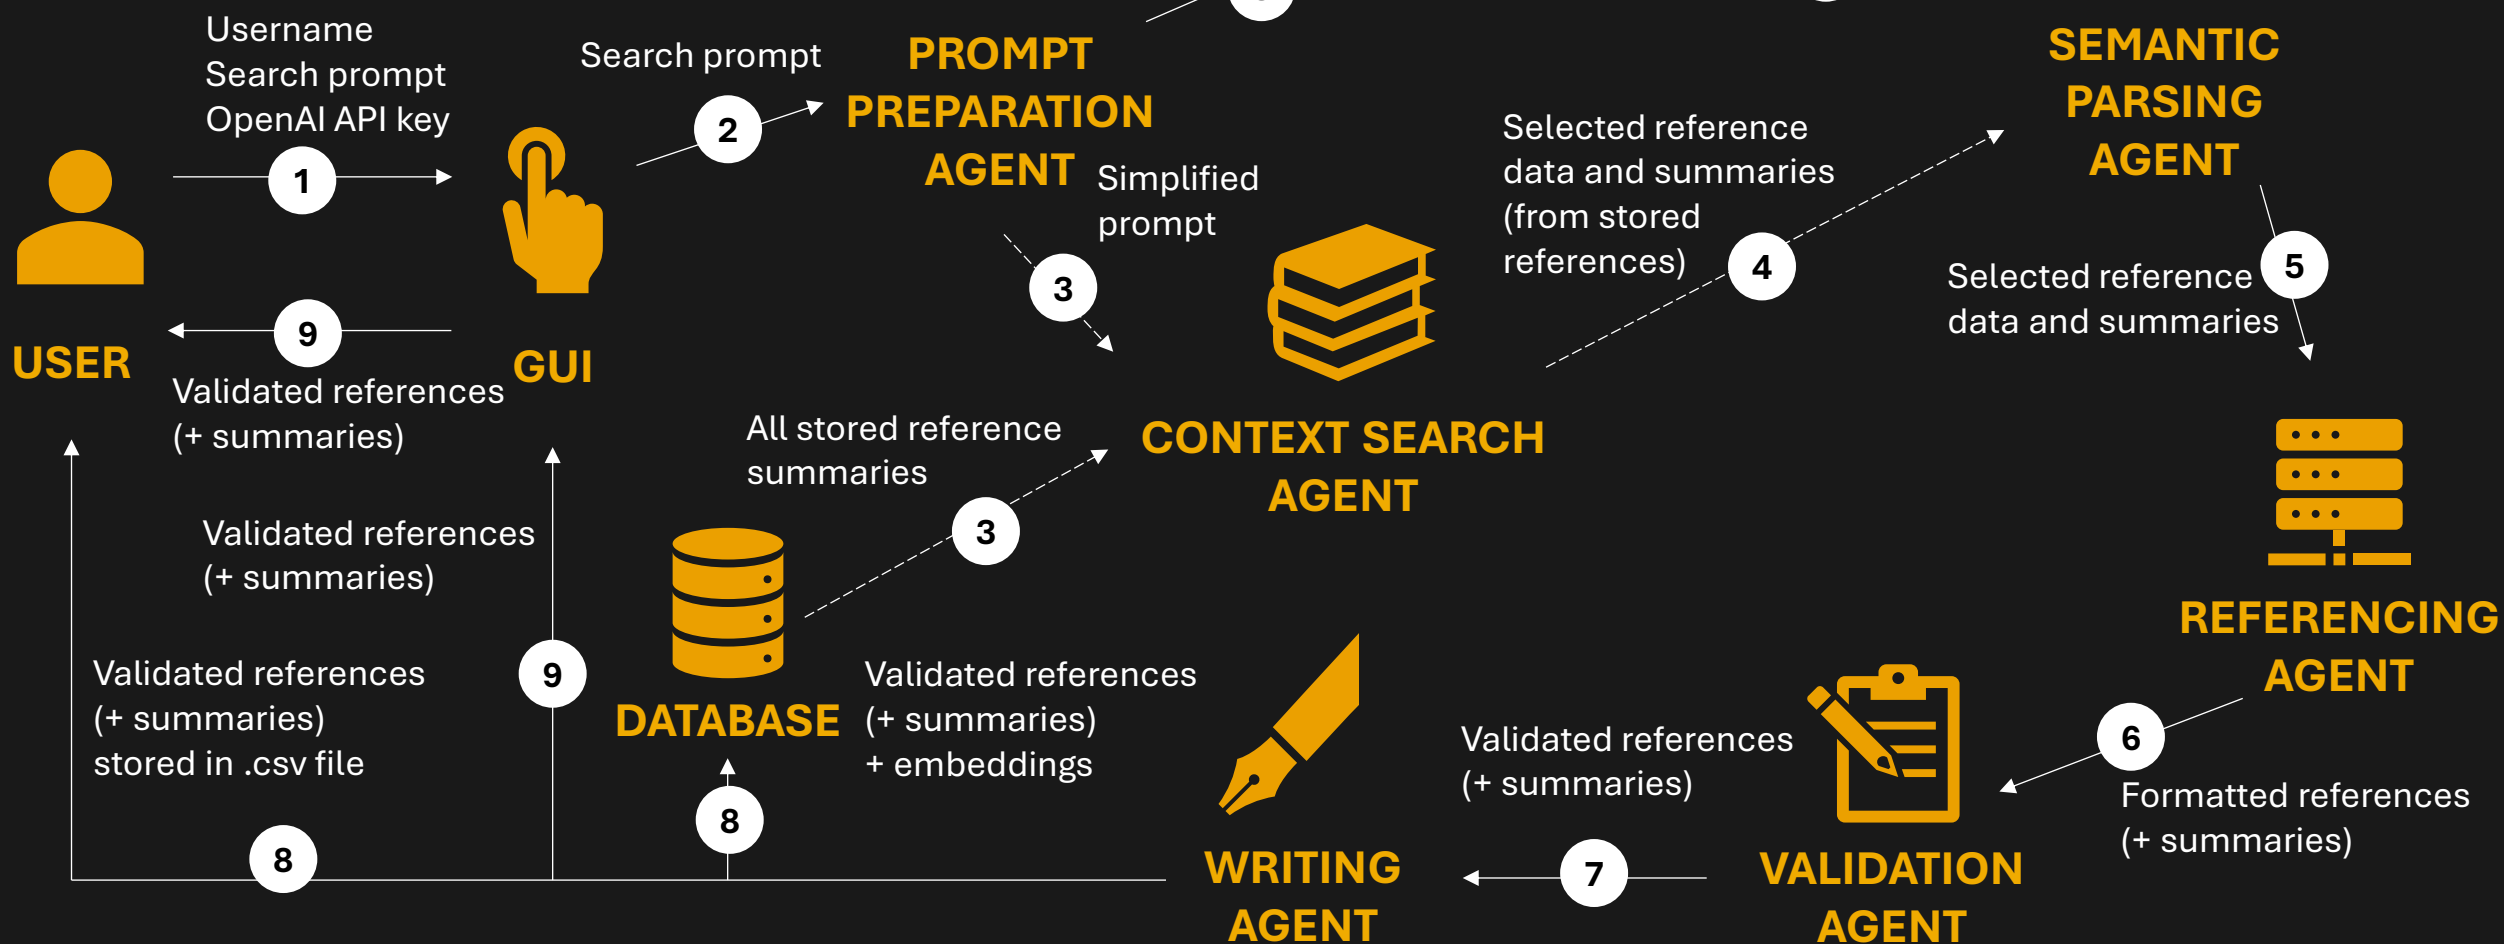
Assignment specifications

1. The agent can **identify and retrieve data**.
2. The **data is processed** in some way.
2. The processed **information is stored/saved/presented**.

Chosen domain: **academic research online**

- **finding results on a website based on search terms** (e.g., social media or a search engine)
- **extracting the data and**
- **sending to an offline location.**

System layout



Application demo

The screenshot displays a Visual Studio Code editor window with a project named "Individual project". The Explorer sidebar on the left shows the project structure, including files like `__pycache__`, `Docs`, `myenv`, `Private`, `OpenAI_api_key.txt`, `SAGE_results`, `intelligent_agents_notebook.ipynb`, `README.md`, `SAGE_database.db`, `SAGE_functions.py`, `SAGE_main.py`, and `SAGE_main.spec`. The Outline sidebar shows a list of agents: General set-up, Load required modules, Create agents (Prompt preparation agent, Search agent, Semantic parsing agent, Referencing agent, Validation agent), Create data storage, Context aware search, Create GUI, Initial username and prompt r..., Return results to user, Compile agent chain, Compile GUI, and Run program.

The main editor area shows the `intelligent_agents_notebook.ipynb` file, which contains the following Python code:

```
user_prompt = input_dict["user_prompt"]

prompt_template = ChatPromptTemplate.from_messages([
    {"role": "system", "content": "You simplify user prompts for academic reference search queries."},
    {"role": "user", "content": "Please simplify the following prompt for use in an academic reference search: {user_prompt}. For example, for an input 'Please provide 5 recent academic references on applications of intelligent agents within management.', you would return: 'intelligent agents in management applications'"}
])

global llm # ensure base model is available for subsequent agents with the same features

llm = ChatOpenAI(
    model="gpt-4o", # most recent OpenAI model, with more advanced reasoning features and more recent context included in training
    database
    temperature=0, # looking for results as accurate as possible and with little room for improvisation or hallucination
    max_tokens=None, # no restriction on prompt size (especially as the model will be parsing lengthy reference inputs at later steps)
    timeout=None,
    max_retries=2,
    api_key=openai_api_key
)

llm_chain = prompt_template | llm | output_parser
simplified_prompt = llm_chain.invoke({"user_prompt": user_prompt})

global response
response = {"user_prompt": user_prompt, "search_prompt": simplified_prompt}
print(f"Prompt preparation completed at {datetime.now()}")
print(f"Output: {response}")
return response
```

The code is executed, and the output is displayed in the terminal:

```
[6] ✓ 0.0s

# use-case testing

test_search_prompt = "Please provide 5 academic references on using natural language processing vs rule-based logic for multi-agent system communication."
# openai_api_key= "REMOVED"
prompt_preparation_agent({"user_prompt": test_search_prompt})

[46] ✓ 1.3s

Prompt preparation completed at 2025-04-12 01:00:13.411464
Output: {'user_prompt': 'Please provide 5 academic references on using natural language processing vs rule-based logic for multi-agent system communication',
'search_prompt': 'natural language processing vs rule-based logic in multi-agent communication'}
```

The terminal output shows the prompt preparation completed at 2025-04-12 01:00:13.411464. The output is a dictionary with 'user_prompt' and 'search_prompt' keys. The 'search_prompt' is 'natural language processing vs rule-based logic in multi-agent communication'.

At the bottom of the screen, there is a "Search agent" section with a text input field containing the text "import requests".



Development issues and troubleshooting

LLM reference search

- ✗ Hallucinations¹
- ✗ Outdated references
- ✗ Unable to circumvent with prompt engineering
- ✓ Add API calls to structured databases (LLM for prompt simplification and reference prioritising)

LLM output parsing

- ✗ Always return string outputs
- ✗ Difficult to ensure correct formatting³
- ✓ Extensive data formatting enforcement (lists/dictionaries)

Referencing

- ✗ Complex referencing rules (depending on citation source and chosen format)
- ✓ Simple and versatile scheme

Reference storage

- ✗ Avoid saving duplicated references
- ▶▶ Removing duplicates before storage (but how best to identify? DOI? combinations?)⁴

Reference validation

- ✗ No clear way to identify problematic data (e.g. pages, authors)²
- ✗ Complex architecture required to prompt user for corrections
- ✓ Simple data type checks and URL validation
- ✓ LLM cross-check (formatting)

Context search

- ✗ Possibly helpful, but difficult to distinguish when old vs new references need to be prioritised
- ▶▶ Allow user to turn context-aware search off

Plus: setting up API connections, managing dependencies, library updates, langchain documentation, etc

¹Maes, 2024; ²Choi et al., 2023; ³Liu et al., 2024; ⁴McKeown and Mir, 2021

Critical reflection

STRENGTHS

- Natural language queries (simplify user interactions)¹
- API calls (overcome LLM hallucination issues, provide solid foundation for LLM parsing)^{1, 2}
- Flexible and scalable database connections
- NLP-based agents (reference prioritisation, flexible formatting implementation)³
- Data validation (structured and flexible scheme)
- LLM-based reference summaries (simplify inspection)
- Retrieval of public data only (no privacy/ethical concerns)⁴
- Local data storage (low cost, high degree of user control)
- Modular code implementation

WEAKNESSES

- Archaic GUI
- Multiple dependencies
- Limited literature database scope
- Suboptimal reference validation scope (and limited capacity to handle and resolve issues)
- NLP-based agents (difficulty handling ambiguous or poorly-defined prompts, LLM bias, possibility of hallucination)^{2, 5}
- Poor scalability and real-time processing capability
- Simple in-memory vector database indexing strategy
- Limited error handling / debugging features



Possible improvements

Area	Improvement
Referencing	Implement a formal reference format (and allow user to select others)
Search prompt	Integrate LLM-generated disambiguation or specification requests to the user
LLM agent	Allow user to edit default prompt instructions, or specify specific LLM of interest
Verification	Integrate user verification / approval before saving references
Validation	Additional reference validation rules or more complex validation steps
Context-aware search	Turn on / off depending on user needs
Database indexing	Develop indexing strategy to avoid saving duplicates
API handling	Batch processing, exponential back-off, usage/monitoring alerts
Error management	Comprehensive exception handling, diagnostic logging, and fallback recovery procedures

Conclusions and learning reflections

- Developed a multi-agent system that successfully:
 - ✓ finds results on (a) website(s) based on search terms (CrossRef, arXiv, Pubmed)
 - ✓ extracts and processes the data (API calls + LLM and rule-based processing)
 - ✓ stores the data in an offline location (structured + vector database, .csv file, and GUI)
- Explored many important and new computer science concepts (data validation and parsing, object-oriented programming, control flow, error handling/debugging/testing, GUI, etc.)
- Experienced advantages and pitfalls of different agent-based architectures (rule-based vs NLP) and challenges of working with multi-agent systems
- Identified strengths, weaknesses, and possible improvements for my proposed implementation
- Wrote my first computer program (with a little help from a few LLMs)

References

- Benji, N. (2025) *LLMs Vs. Deterministic Logic — Overcoming Rule-Based Evaluation Challenges*, Medium. Available at: <https://blog.gopenai.com/llms-vs-deterministic-logic-overcoming-rule-based-evaluation-challenges-8c5fb7e8fe46> (Accessed: 13 April 2025).
- Brynjolfsson, E. and McAfee, A. (2014) *The Second Machine Age: Work, Progress, and Prosperity in a Time of Brilliant Technologies*. 1st edn. W. W. Norton & Company.
- Cao, L. (2021) 'AI in Finance: Challenges, Techniques and Opportunities'. arXiv. Available at: <https://doi.org/10.48550/arXiv.2107.09051>.
- Choi, W. *et al.* (2023) 'Building an annotated corpus for automatic metadata extraction from multilingual journal article references', *PLOS ONE*, 18(1), p. e0280637. Available at: <https://doi.org/10.1371/journal.pone.0280637>.
- Chui, M., Manyika, J. and Miremadi, M. (2016) *Where machines could replace humans—and where they can't (yet) | McKinsey, McKinsey Digital*. Available at: <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/where-machines-could-replace-humans-and-where-they-cant-yet> (Accessed: 13 April 2025).
- *Data protection explained - European Commission* (no date). Available at: https://commission.europa.eu/law/law-topic/data-protection/data-protection-explained_en (Accessed: 13 April 2025).
- Juziuk, J., Weyns, D. and Holvoet, T. (2014) 'Design Patterns for Multi-agent Systems: A Systematic Literature Review', in O. Shehory and A. Sturm (eds) *Agent-Oriented Software Engineering: Reflections on Architectures, Methodologies, Languages, and Frameworks*. Berlin, Heidelberg: Springer, pp. 79–99. Available at: https://doi.org/10.1007/978-3-642-54432-3_5.
- Lalwani, A. *et al.* (2025) 'Autoformalizing Natural Language to First-Order Logic: A Case Study in Logical Fallacy Detection'. arXiv. Available at: <https://doi.org/10.48550/arXiv.2405.02318>.
- Liu, Y. *et al.* (2024) 'Are LLMs good at structured outputs? A benchmark for evaluating structured output capabilities in LLMs', *Information Processing & Management*, 61(5), p. 103809. Available at: <https://doi.org/10.1016/j.ipm.2024.103809>.
- Maes, S. (2024) 'Fixing Reference Hallucinations of LLMs'. Available at: <https://zenodo.org/records/14791389> (Accessed: 13 April 2025).
- McKeown, S. and Mir, Z.M. (2021) 'Considerations for conducting systematic reviews: evaluating the performance of different methods for de-duplicating references', *Systematic Reviews*, 10(1), p. 38. Available at: <https://doi.org/10.1186/s13643-021-01583-y>.
- Russel, S. and Norvig, P. (2021) *Artificial intelligence: a modern approach*. 4th edn. Upper Saddle River, NJ : Prentice Hall: Pearson. Available at: <https://doi.org/10.1109/MSP.2017.2765202>.
- Salau, L. *et al.* (2022) 'State-of-the-Art Survey on Deep Learning-Based Recommender Systems for E-Learning', *Applied Sciences*, 12(23), p. 11996. Available at: <https://doi.org/10.3390/app122311996>.
- Topol, E. (2019) *Deep Medicine: How Artificial Intelligence Can Make Healthcare Human Again*. 1st edn. NY: Basic Books. Available at: <https://psnet.ahrq.gov/issue/deep-medicine-how-artificial-intelligence-can-make-healthcare-human-again> (Accessed: 13 April 2025).
- Woolridge, M. (2009) *An Introduction to MultiAgent Systems, 2nd Edition | Wiley*. 2nd edn. Chichester: John Wiley & Sons. Available at: <https://www.wiley.com/en-us/An+Introduction+to+MultiAgent+Systems%2C+2nd+Edition-p-9780470519462> (Accessed: 11 February 2025).
- Yadav, A., Patel, A. and Shah, M. (2021) 'A comprehensive review on resolving ambiguities in natural language processing', *AI Open*, 2, pp. 85–92. Available at: <https://doi.org/10.1016/j.aiopen.2021.05.001>.